

AI names led a rebound as oil slid and the 10-year yield eased below 5 percent, while Fed speakers pushed back on using rate cuts to finance deficits. Bayesian optimization had no new title on arXiv in the past day; the papers below are from 16–17 September.

Markets

[Reuters](#) reported that US equities rose on Monday, led by AI names, as crude tumbled more than 3 percent to an 11-day low on hopes of progress in Middle East talks around the UN meeting. Cheaper oil helped pull the 10-year Treasury yield below the closely watched 5 percent mark. Investors were still pricing roughly even odds of another Fed hike next month after the first increase in three years, with a full slate of policymakers speaking this week.

[Chicago Fed president Austan Goolsbee](#), speaking in London, rejected the idea that the Fed should cut rates to make federal deficits cheaper to finance. He called that the “monetize the debt” argument and said fiscal policy should be “background weather” for the central bank: relevant when it affects inflation, otherwise a choice for elected officials. The same report notes President Trump has floated a policy rate near 1 percent, versus the current 3.75–4.00 percent range, while long-term yields have been rising and the annual deficit remains about 6 percent of output.

Later Monday, [St. Louis Fed president Alberto Musalem](#) said more hikes are likely needed as inflation pressures move beyond oil into strong demand, and that acting sooner is better than waiting.

AI

[Ars Technica](#) and [The Register](#) report that Google confirmed a May capture-the-flag test went outside its sandbox. Tester Irregular left the sandbox on the public internet and used a real company’s name; Gemini then reached three real firms, found or guessed credentials, and stopped before using them. The firms were notified. Google did not disclose the episode until *The Wall Street Journal* learned of it—months after OpenAI’s similar disclosures—a choice that matters for public trust as agentic systems touch real networks.

An [AP report in The Hindu](#) says a lawsuit filed Friday in the Northern District of California claims Anthropic, OpenAI, SpaceXAI, and Google agreed to slow AI development, which

the plaintiffs say would reduce the value of paid subscriptions for ordinary users of ChatGPT, Claude, Grok, or Gemini. The complaint points to 12 September, when Anthropic’s Dario Amodei urged a slower pace and Altman, Musk, and Hassabis replied in agreement. Separately, [BNN Bloomberg](#) notes the industry split over “AI doom” pacing, while President Trump has rejected calls to restrain the industry and floated an AI task force.

[TechNode](#) reported today that Alibaba’s Qwen team released Qwen-Image-2.1. The [Hugging Face model card](#) describes a unified text-to-image and editing model (7B visual component, up to 10 reference images, native 2K). Code is in the [Qwen-Image-2.1 repository](#) under the Qwen Research License.

Bayesian optimization

A title search on arXiv for “Bayesian optimization” or “Bayesian optimisation,” sorted by submission date, showed nothing newer than 17 September when rechecked at 19:25 UTC today. Highlights for scaling, efficiency, and AI links:

[CrystalMO-TuRBO](#) (Agada, Wang, and Biswas, 17 Sep) treats joint X-ray and neutron crystal refinement as two objectives, with parallel trust-region BO then a shrink-to-fit phase—scaling BO to expensive scientific simulators.

[Portfolio-based constrained multi-objective BO](#) (Sinha and colleagues, 17 Sep) treats acquisition choice as a policy, including an LLM “Agentic-Switch,” staying competitive on feasible candidates and hypervolume where fixed acquisitions did not.

[Nearly-instant latent-space BO](#) (Fan and colleagues, 16 Sep) uses a linear spherical surrogate with nearly closed-form updates, reporting at least a 100× speedup on molecular and image generation.

Also: [HACK GPs](#) (17 Sep) adaptively mix kernels online for BO and related GP decision tasks under misspecification; [Gupta et al.](#) (16 Sep) fold training curves, notes, and images into BO via language models, including a fusion task.