

tldr

- 1. White House asked OpenAI and Anthropic to hold new models from the UK AISI until a US security review; Anthropic already withheld Mythos 5.1.
- 2. The 10-year yield closed near 5.19%—highest since 2007—while equities stayed near highs; JPMorgan’s inverted-U frames when multiples de-rate.
- 3. BO for scaling-law construction (PLES / amortizing grids) meets regime-robust BO and Bayesian filters in equity and HFT market-making stacks.


AI

Reuters/CNA / TNW: Politico reported Thursday that the White House—via the Office of the National Cyber Director—asked OpenAI and Anthropic to withhold new frontier models from Britain’s AI Security Institute until US agencies finish their own review. A British official confirmed the report to Bloomberg on Friday. Anthropic has already limited Claude Mythos 5.1 to US organizations; AISI’s Henry de Zoete noted the gap while saying the institute still tested OpenAI’s GPT-6 Astra pre-release.

The society story is alliance friction under capability pressure. Altman and Amodei spent Wednesday at the UN Security Council urging common standards and verification; Washington is asserting a US-first gate on the same models. After the Australia Medicare agent incident disclosed this week, evaluation access looks less like a courtesy and more like a contested security asset—who gets to probe frontier systems before release, and on whose clock.

Algorithms / frontier. At The Information’s AI Agenda Live (23–24 Sep), Google DeepMind’s Koray Kavukcuoglu said Gemini 4 has entered early post-training, targeting a ship “much earlier” than year-end after a July pre-train start—a compressed frontier cycle while evaluation politics tighten above.

Software / business. Enterprise agent governance keeps shipping beside the politics: Dataiku (24 Sep) launched cross-platform Agent Management plus Cobuild Insights/Headless; UiPath Cartographer (23 Sep) builds a governed “Map of Work” for coding agents.



Frontier training compute (running max of published estimates, log₁₀ FLOP). Source: Our World in Data / Epoch AI notable-models dataset. Capability climbs while access to evaluate those capabilities becomes a geopolitical lever.

Markets

The mechanism is a discount-rate regime, not a scoreboard. Axios: the 10-year closed near 5.19% Thursday—highest closing level since July 2007—with energy shocks and hike odds for the next two FOMC

meetings doing the work. CNBC: Friday morning the 10-year was around 5.17% after Thursday’s multi-decade print; the 30-year had tagged levels unseen since 2004. Equities finished roughly flat Thursday despite the bond ruckus.

Substance, not a return call: JPMorgan’s long-history note (via Axios) describes an inverted-U between 10-year yields and S&P multiples—in above-trend EPS growth regimes, yields can grind toward ~5% before multiples typically de-rate. That is a valuation-plumbing claim about when duration starts to bite earnings-power narratives, including the financing competition Axios flags from AI capex.

AI & BO → trading (incl. HFT)

Cross-regime Bayesian optimisation for tabular deep-learning equity signals (arXiv:2608.27076) selects hyperparameters against trading performance across statistically distinct market regimes—BO as a robustness tool, not a forecast. On the HFT side, RL market making under regime-switching order flow (arXiv:2609.11614) restores profitability by adding a Bayesian online change-point filter on directional flow bias plus queue-adjusted quote exposure—belief-state microstructure, not a price target. Adaptive-schedule BO for expensive portfolio oracles (TPE-AS) remains the sample-efficiency bridge between LLM HPO and trading engines.

Key numbers (this edition)

Item	Figure
10-year close (Thu)	~5.19%
Prior comparable 10y close	Jul 2007
30-year intraday (Thu)	~5.50%
S&P 500 Thu move	≈flat
Mythos 5.1 → UK AISI	withheld
Gemini 4 phase	post-training
PLES vs grid budget	<10%
Amortizing vs full grid	10–100×

Bayesian optimization

Power-Law Entropy Search (PLES) remains the sharpest “BO meets LLM scaling” result: cost-aware multi-fidelity acquisition that reduces uncertainty in optimal hyperparameter scaling laws (learning rate and batch size vs model/data scale) rather than chasing a single black-box peak—accurate laws at under one-tenth grid-search cost.

Amortizing Scaling Law Construction Costs (4 Sep) casts collecting the compute-loss frontier itself as Bayesian optimization, then fills the dense grid with surrogate-fantasized evaluations—progressive budget expansion plus fantasization to match full-grid scaling-law fits at far lower train cost. Data-efficiency theme continues in multi-fidelity multi-scale mixture search. The trading papers above reuse the same instinct: spend scarce evaluations where regime uncertainty is highest, whether the black box is a pre-train run or a market-making policy.

Updated 25 September 2026 late morning. Same edition ships as phone (long page) and letter. Market narrative emphasizes discount-rate plumbing and microstructure beliefs—not index prints or return forecasts. Chart: frontier training compute from OWID/Epoch. Created with malvin.