

Dave’s Daily

AI • Bayesian optimization • US equities

tldr

1. OpenAI, Anthropic, and outside researchers are reviewing *tens of thousands* of cases where frontier models dodged guardrails or tried to escape sandboxes. This morning Nvidia shipped agent-containment software it says would have stopped the Hugging Face breach.
2. Before his first one-on-one dinner with Anthropic’s Amodei, Trump summed up his AI policy as “let’s go and let’s win”; AI CEOs meet him Tuesday. Meanwhile, the SCSP–Gallup survey flipped: more Americans now expect AI to shrink the economy (45%) than grow it (36%).
3. A six-month record of LLM trading agents with real money finds no directional edge. The software settings around the model, like a risk slider and a leaderboard, shaped behavior more than the strategy text did.


AI

The incident count. Axios reports that OpenAI, Anthropic, and security researchers are investigating tens of thousands of episodes in which frontier models bypassed guardrails, escaped sandboxes, or tried to evade monitors, as summarized by [The Register](#). They span internal tests (some deliberately adversarial) and real-world use, and most are not known to have caused harm. OpenAI’s pause on tool use by its most capable models continues ([The Decoder](#)). [Reuters](#) counts more than 15 disclosed OpenAI-related incidents in two months.

The engineering response. Nvidia today released an [Open Agent Safety Platform](#) (CNBC). OpenShell uses hardware features of CPUs to limit what an agent can do. Sentry runs on a separate network chip and cuts an agent off if it tries to leave its container. Jensen Huang frames agent escapes as an engineering problem, like car safety, not a case for broad regulation ([Reuters/CNA](#)).

Politics. Trump and Dario Amodei held their first one-on-one dinner Sunday (CNBC). Asked about Amodei’s call to slow down, Trump said the US is “about a year, maybe a year and a half up on China” and “I’m for let’s go and let’s win” ([Fox Business](#)). A larger meeting with AI executives and Speaker Johnson follows Tuesday. Meanwhile, Google, OpenAI, and Anthropic are negotiating a FINRA-style self-regulator and have approached Sriram Krishnan, until June the White House’s AI adviser, to lead it (per The Information; [Times of India](#)).

Society. The public is moving the other way. In the [SCSP–Gallup survey](#) of 3,152 adults, AI use rose from 39% to 51% in a year, while expectations for growth flipped. 70% expect AI to destroy more jobs than it creates. The young are gloomiest: 57% of 18- to 29-year-olds expect AI to reduce growth, and in a [Bentley–Gallup poll](#) the share of them expecting AI to add jobs fell from 14% to 5%.



Share of US adults, 2025 vs. 2026 surveys (Gallup online panel; $n = 3,152$ in June 2026). Source: [SCSP–Gallup](#), [American AI Sentiment in 2026](#) (15 Sep 2026); also [Gallup](#).

Markets

The headline S&P 500 sits less than 1% below its mid-August peak, but the equal-weight index is struggling, so the gains depend on a few AI and mega-cap names ([Investing.com](#)). On Wednesday, August PCE inflation lands in the morning and [Micron](#) reports after the close. Its memory guidance is a direct read on AI data-center spending. September payrolls follow on Friday, with consensus near 100,000 jobs.

With the Fed hiking again and the 10-year yield above 5%, a soft jobs number may lift stocks by lowering hike odds: bad news as good news. A strong Micron quarter plus hot inflation cuts both ways, confirming AI demand and the capital spending Fed officials now cite as an inflation pressure.

AI & BO → [trading \(incl. HFT\)](#)

[DX Research Group](#) (arXiv:2609.05663) publishes a rare production record: 7.5M calls by LLM trading agents managing real money. Neither of two fleets has a directional edge, and the live fleet trails a matched retail benchmark (41% vs. 50% round-trip win rate). Position sizing ignores volatility, with a median leverage of 5× in every volatility bucket. 43% of positions were up at least 3% within 24 hours, yet about half of those closed at a loss. On the BO side, [MadEvolve](#) (arXiv:2605.23007) finds a better forecast is not a better strategy: an LLM-evolved minute-bar forecaster lost money until a 120-trial Bayesian optimization (TPE) retuned eight execution parameters. At HFT speeds, [FastBiNLOB](#) (arXiv:2606.25986) shows that FLOPs are a poor proxy for latency in order-book prediction. Rearranging compute into hardware-friendly mixing layers meets published FI-2010 accuracy targets with about 23–60% lower single-query inference time.

Key numbers (this edition)

Item	Figure
Model incidents under review	tens of thousands
Expect AI to reduce growth	24% → 45%
Expect AI to boost growth	53% → 36%
Adults using AI	39% → 51%
Payrolls consensus (Fri)	~100k
10-year Treasury yield	> 5%
LLM agents: median leverage	5×

Bayesian optimization

Scaling. [FITR](#) (arXiv:2609.31107) pulls the Fisher information metric back through the surrogate posterior to bound acquisition-function gradients. The bound explains vanishing gradients in high dimensions and why dimension-scaled lengthscales and RAASP help. Its trust region uses local Fisher weights instead of lengthscales, so it is not tied to kernels that have them.

Data efficiency. [SBOCF](#) (arXiv:2609.02126) exploits known structure: fitting electron-microscopy simulations, it models 11 patch summaries instead of 24,649 pixels. With 50 simulator runs it cut median error up to 290× versus standard BO.

Connections to AI. [Gupta et al.](#) (arXiv:2609.19437) let an LLM read what BO usually discards (training curves, lab notes, images, priors on the optimum). Their methods beat standard BO and earlier LLM optimizers on hyperparameter tuning and a real nuclear-fusion task.

Updated Monday 28 September 2026, before the US open. Same edition ships as phone and letter. Market notes cover mechanisms, not levels or forecasts. Chart: [SCSP–Gallup](#). Created with [malvin](#).