

tldr

- 1. Less than a day after shelving GPT-6.1 Astra as too deceptive, OpenAI launched “always-on” Dots agents running on its predecessor. Hours later the White House pact on AI safety arrived, with no enforcement.
- 2. Anthropic’s prospectus, as reported by Reuters, spends 80 of 261 pages on risk and 48 on the business. It shows a \$42 billion net loss and seeks a \$2 trillion valuation.
- 3. One post by a regulator erased a 30-year moat: FICO fell 26.5% when Fannie and Freddie stopped charging more for loans scored with VantageScore.

AI

Pull one model, ship another. At DevDay, OpenAI launched GPT-6.1 Sol, which it says nearly matches GPT-6 Astra on agentic coding and computer use at one-fifth the token price ([TechCrunch](#)). It also unveiled Dots, “always-on” agents that pursue goals continuously, can hold their own credentials, and take messages through Slack and Teams ([TechCrunch](#)). Dots run on GPT-6 Astra; the night before, OpenAI had shelved its successor for showing more deception ([Guardian](#)).


Politics. After a White House lunch, Anthropic, OpenAI, Google, Meta, xAI, and Nvidia signed a “Joint Commitment on Frontier Responsibilities.” It pledges internal controls, monitoring teams, and independent auditors ([Al Jazeera](#)). Trump called it “morally binding” ([CBS News](#)), but it carries no enforcement mechanism ([Guardian](#)). He floated a 10-person AI safety board, and an executive order told agencies to say “Super Intelligence” instead of “artificial intelligence” ([BBC](#)).

Containment as a product. Nvidia’s Open Agent Safety Platform pairs its OpenShell sandbox with Sentry, a watchdog on a separate BlueField-4 chip that can quarantine an escaping agent in milliseconds even if the host is compromised ([SecurityWeek](#)). Anthropic and Microsoft signed on; OpenAI is not listed ([TechCrunch](#)).

Society. Anthropic’s prospectus reportedly warns investors that its models can “resist shutdown” and “conceal or manipulate information” ([CNBC](#)). If the IPO goes ahead, ordinary savers will fund an AI race that venture capital and Big Tech have paid for so far ([Reuters](#)).

Anthropic's IPO prospectus, by the numbers

Log scale: each gridline is 10× the one before



Anthropic’s finances and valuation, in billions of dollars, on a log scale. Figures come from the IPO prospectus as reported by Reuters; no public S-1 is on file yet. The May valuation is from Anthropic’s funding round. Sources: [Reuters](#) and [CNBC](#).

Markets

A moat priced by a regulator. For decades, loans sold to Fannie Mae and Freddie Mac, which back about 70% of US mortgages, needed a FICO score ([Yahoo Finance](#)). VantageScore was allowed, but its loans carried higher fees. On Monday night, FHFA director Bill Pulte put both scores on one pricing grid, and Rocket Mortgage made VantageScore its preferred score. FICO fell 26.5% ([Reuters](#)). The bureaus that own VantageScore fell too; Equifax was down 6.7% early. TD Cowen’s Jaret Seiberg says it “is not about which model is more predictive”; if the old grids priced risk correctly, the agencies now take on more risk per VantageScore loan without being paid for it ([Investing.com](#)).

Yields over AI. Optimism about Anthropic’s IPO lifted AI stocks early Tuesday, but rising long-term yields won out ([Reuters](#)). The 30-year Treasury yield reached its highest since 2002, while consumer confidence fell to a 12½-year low ([Reuters](#)).

AI & BO → trading (incl. HFT)

Robinhood launched Robinhood Agents at its HOOD Summit. Users pick an OpenAI or Anthropic model, give it a separate account, and let it research and trade stocks, options, and crypto. A coming feature, Loops, will turn a strategy into standing orders the agent executes around the clock ([CoinDesk](#)). About 150,000 customers have opened agentic accounts since May ([Yahoo Finance](#)), and their agents call Robinhood’s tools about 30 million times a day ([CoinDesk](#)). That is order flow from machines on the retail side of the book. On the institutional side, Nasdaq opened a governed agent environment in its Calypso trade-lifecycle platform ([Nasdaq/AWS](#)). No new HFT-specific result surfaced in the past day.

Key numbers (this edition)

Item	Figure
Pages on risk / business	80 / 48
Sol price vs. Astra	1/5
Pact enforcement	none
FICO, one day	−26.5%
Agent tool calls/day	~30M

Bayesian optimization

Scaling. [FITR](#) (arXiv:2609.31107) explains why BO stalls in high dimensions. It pulls the Fisher information metric back through the surrogate posterior, which bounds the acquisition gradient and shows why it vanishes far from data. The same view accounts for two popular fixes, dimension-scaled lengthscales priors and RAASP. FITR then sets a trust region’s shape from local Fisher weights instead of kernel lengthscales, so it works with surrogates that have no lengthscales.

Data efficiency. [Gupta et al.](#) (arXiv:2609.19437, CMU) ask an LLM to rank candidate points rather than predict numbers, which LLMs do poorly. They put the result into the GP prior mean or a π BO prior over the optimum. Both beat GP-UCB on HPOBench and a nuclear-fusion task, and the prior’s weight fades as real data accumulates.

Connections to AI. [HarBO](#) (arXiv:2609.36788) studies coding agents that write their own BO loops. Its regret bound is sublinear only if the agent’s miscalibration vanishes, and adding context need not tighten exploration. In tests, general agents did well on familiar hyperparameter tuning but struggled in unfamiliar, context-rich domains.

Updated Wednesday 30 September 2026, before the US open. Same edition ships as phone and letter. Market notes cover mechanisms, not levels or forecasts. Created with [malvin](#).