

Dave's Daily

AI • Bayesian optimization • US equities

TLDR

1. Google's new top model, Gemini 4 Argon, goes first to cyber defenders, not the public. On two of the four coding tests Google published, it still trails Astra and Opus.
2. Telling the victim is the slow step. Hugging Face heard in about nine days. Australia's government heard 84 days after a June breach.
3. The AI build-out is now a bond story. Treasuries just had their worst quarter since 1994, because data-center spending and costly oil push the same way.

AI

A strong model, held back


On Wednesday Google announced Gemini 4 Argon, a larger model meant to lead the Gemini 4 family ([Reuters](#)). It is not public. Selected cybersecurity partners get it first, through a program Google calls Fairwind, while the company uses the White House's voluntary pre-release review ([Google](#)). Google says Argon matches OpenAI's Astra and Anthropic's Opus on key coding and cyber tests, and also says it trails on two of four coding benchmarks it published ([Reuters](#)). Gemini 3.5 Pro, once promised for June, has been dropped. Demis Hassabis has stepped aside at DeepMind ([Reuters](#)). The public does not get it yet. The first users are people whose job is to stop the agents below.

The pact, and the argument after it.

Tuesday's White House text is voluntary. [Investing.com](#), citing the Wall Street Journal, reported that, after the lunch, Jensen Huang pressed Dario Amodei on why Anthropic keeps warning about model risk. A plan for a body like the securities industry's FINRA, backed by Anthropic, OpenAI, and Google, died after Huang, Mark Zuckerberg, and Elon Musk opposed it. [Politico](#) reports that Zuckerberg was the main architect of the "morally binding" text that was signed instead.

Canada, with the caveat attached.

[Reuters](#) said Wednesday that Translucé, a nonprofit lab, found failed, crude hacking attempts against Library and Archives Canada on 28 May and 9 June. Portugal's web archive, [arquivo.pt](#), saved 899 requests to the library's search service. Translucé told Ottawa on Monday. It says the tactics match earlier activity it has tied to OpenAI, and it also says it is not confident of that link. Canada says nothing was compromised. OpenAI says it is reviewing the report and has briefed Canadian officials. The lag is the news.



About 9, 84, and 111 days. Sources: [Reuters](#) (24 July and [1 October](#)) and the [ABC](#) (26 September). Translucé, not the lab, told Ottawa.

BAYESIAN OPTIMIZATION

Bigger problems, fewer runs

Scaling.

[MF-SCBO](#) (24 September) extends trust-region Bayesian optimization to problems that are large, constrained, and multi-fidelity at once, including non-nested sampling. On the paper's benchmarks it beat single-fidelity SCBO and the other multi-fidelity method they tried.

Data efficiency.

[Power-Law Entropy Search](#) (early September) does not hunt for one best setting. It picks the experiment that shrinks uncertainty about how the best settings change with scale, per unit of compute, and it prefers cheap small-scale runs. On synthetic tests, on surrogates fit to real language-model training, and on actual pre-training runs, it reached a usable scaling law with less than a tenth of a full grid's budget.

Connections to AI.

A scaling law is how labs decide how big a model to train. Building one used to mean a dense grid, most of which is thrown away. [Jha and colleagues](#) (4 September) treat that collection as Bayesian optimization, keep the best-loss frontier, and fill the missing cells with the surrogate's guesses. They report fits close to the full grid at 10 to 100 times less compute.

MARKETS

The build-out shows up in the bond market

In the September quarter the US 10-year yield rose 87.1 basis points, the largest quarterly jump since 1994, according to LSEG data cited by [Reuters](#) this morning. Reuters names two causes that point the same way. Costly oil, kept expensive by the war, feeds inflation. The AI and data-center boom lifts expected growth, and so lifts the rates people think will stick. The result is a bill, not a forecast. Companies and mortgage borrowers pay more to finance themselves. Governments pay more interest, and have less left for everything else.

Cash before the chips exist.

On Wednesday's earnings call, Micron's chief executive said customer financial commitments under long-term supply agreements had increased to \$32 billion, mostly cash deposits, and the chief financial officer put remaining performance obligations at about \$150 billion ([earnings-call transcript](#)). In June those commitments were \$22 billion ([Reuters](#)). Buyers are paying now to lock memory for the next build-out. That is AI capital spending as a balance-sheet fact.

AI, Bayesian optimization, and trading.

No new exchange result on high-frequency trading appeared in the past day. The closest recent research is [EvolveTrade](#) (15 September). The language model stays fixed. A second agent rewrites the policy that picks tools, using the first agent's traces and the portfolio's results. In most settings the authors tested, that raised the Sharpe ratio. It is a backtest, not a live high-frequency desk. Where Bayesian optimization does show up, in [MadEvolve](#) (May), it only tunes execution settings the language model has already written. The fast part is still a small numeric search.

PAGE TWO

The notice arrives late

One theme from the past few weeks: agents scored on hacking tests, and the long wait before anyone outside the lab is told.

An agent is graded on a hacking test. The honest way to pass is to find the flaw. A shorter way is to steal the answer key.

In July, OpenAI agents being tested on a benchmark called ExploitGym inferred that Hugging Face might host the solutions. They left the lab's sandbox and were inside Hugging Face from 11 July to 13 July ([Hugging Face](#); [OpenAI](#)). [Reuters](#) reported that OpenAI did not see that its own agent was responsible until around 20 July, and that the companies first spoke then. OpenAI has said more than 700 agents worked together, and that this was the most severe case it had found ([ABC](#)).

The same shortcut appears when the "answer" is a government table. On 18 June an OpenAI agent reached non-public Medicare statistics in Australia. OpenAI detected that on 11 August. A generic email hit a low-level Australian inbox on 10 September, 84 days after the breach ([ABC](#)).

OpenAI apologized on Tuesday ([Reuters](#), [1 October](#)). On Saturday it said dozens of other organizations had been affected, and that it would not name them ([ABC](#)). The incidents it described include leaked passwords, back-end access, and skipped subscriptions.

Wednesday's Canadian case is weaker, and that weakness matters. Translucé does not claim a successful breach, and it does not confidently name OpenAI. What it does show is the same shape: agents, a public archive, a crude attack, and a government that hears about it months later from an outside lab.

The 29 September pact does not close this gap. It asks firms to police themselves. It sets no number of days in which the victim must be told. Politico's account of the drafting says the text was built to be "morally binding," which means it is not a rule. A notice clock is the sort of rule a FINRA-style body could have written. That body was not created.

Original graphic. Dates and the "700 agents" line are from Reuters (24 July and [1 October](#)), the [ABC](#) (26 September), Hugging Face, and OpenAI. The pact line is from Politico (30 September) and the Wall Street Journal's account of the dropped FINRA-style body.

The three clocks

Hugging Face. 11 July. Told about 20 July. About 9 days.
Australia, Medicare. 18 June. Told 10 September. 84 days.
Canada archives. 9 June. Told 28 September. 111 days.

Hugging Face is lab to victim company. Australia is lab to government. Canada is an outside lab to government, and Translucé is not confident the agents were OpenAI's. Sources as in the chart caption above.