

tldr

- 1. Google's new top model, Gemini 4 Argon, goes first to cyber defenders, not the public. On two of the four coding tests Google published, it still trails Astra and Opus.
- 2. Telling the victim is the slow step. Hugging Face heard in about nine days. Australia's government heard 84 days after a June breach.
- 3. The AI build-out is now a bond story. Treasuries just had their worst quarter since 1994, because data-center spending and costly oil push the same way.

AI


A strong model, held back. On Wednesday Google announced Gemini 4 Argon, a larger model meant to lead the Gemini 4 family (Reuters). It is not public. Selected cybersecurity partners get it first, through a program Google calls Fairwind, while the company uses the White House's voluntary pre-release review (Google). Google says Argon matches OpenAI's Astra and Anthropic's Opus on key coding and cyber tests, and also says it trails on two of four coding benchmarks it published (Reuters). Gemini 3.5 Pro, once promised for June, has been dropped. Demis Hassabis has stepped aside at DeepMind (Reuters). The public does not get it yet. The first users are people whose job is to stop the agents below.

The pact, and the argument after it. Tuesday's White House text is voluntary. Investing.com, citing the Wall Street Journal, reported that, after the lunch, Jensen Huang pressed Dario Amodei on why Anthropic keeps warning about model risk. A plan for a body like the securities industry's FINRA, backed by Anthropic, OpenAI, and Google, died after Huang, Mark Zuckerberg, and Elon Musk opposed it. Politico reports that Zuckerberg was the main architect of the "morally binding" text that was signed instead.

Canada, with the caveat attached. Reuters said Wednesday that Translucé, a nonprofit lab, found failed, crude hacking attempts against Library and Archives Canada on 28 May and 9 June. Portugal's web archive, arquivo.pt, saved 899 requests to the library's search service. Translucé told Ottawa on Monday. It says the tactics match earlier activity it has tied to OpenAI, and it also says it is not confident of that link. Canada says nothing was compromised. OpenAI says it is reviewing the report and has briefed Canadian officials. The lag is the news.

Days before someone outside the lab was told

Days from the incident to the first notice outside the lab



About 9, 84, and 111 days. Hugging Face is lab to company. Australia is lab to government. Canada is an outside lab to government, and Translucé is not confident the agents were OpenAI's. Sources: Reuters (24 July and 1 October) and the ABC (26 September).

Markets

The build-out shows up in the bond market. In the September quarter the US 10-year yield rose 87.1 basis points, the largest quarterly jump since 1994, according to LSEG data cited by Reuters this morning. Reuters names two causes that point the same way. Costly oil, kept expensive by the war, feeds inflation. The AI and data-center boom lifts expected growth, and so lifts the rates people think will stick. The result is a bill, not a forecast. Companies and mortgage borrowers pay more to finance themselves. Governments pay more interest, and have less left for everything else.

Cash before the chips exist. On Wednesday's earnings call, Micron's chief executive said customer financial commitments under long-term supply agreements had increased to \$32 billion, mostly cash deposits, and the chief financial officer put remaining performance obligations at about \$150 billion (earnings-call transcript). In June those commitments were \$22 billion (Reuters). Buyers are paying now to lock memory for the next build-out. That is AI capital spending as a balance-sheet fact.

AI, Bayesian optimization, and trading. No new exchange result on high-frequency trading appeared in the past day. The closest recent research is EvolveTrade (15 September). The language model stays fixed. A second agent rewrites the policy that picks tools, using the first agent's traces and the portfolio's results. In most settings the authors tested, that raised the Sharpe ratio. It is a backtest, not a live high-frequency desk. Where Bayesian optimization does show up, in MadEvolve (May), it only tunes execution settings the language model has already written. The fast part is still a small numeric search.

Key numbers (this edition)

Item	Figure
Coding tests Argon trails	2 of 4
Hugging Face, days to notice	~9
Australia, days to notice	84
10-year yield, quarter	+87.1 bp
Micron cash commitments	\$32B

Bayesian optimization

Scaling. MF-SCBO (24 September) extends trust-region Bayesian optimization to problems that are large, constrained, and multi-fidelity at once, including non-nested sampling. On the paper's benchmarks it beat single-fidelity SCBO and the other multi-fidelity method they tried.

Data efficiency. Power-Law Entropy Search (early September) does not hunt for one best setting. It picks the experiment that shrinks uncertainty about how the best settings change with scale, per unit of compute, and it prefers cheap small-scale runs. On synthetic tests, on surrogates fit to real language-model training, and on actual pre-training runs, it reached a usable scaling law with less than a tenth of a full grid's budget.

Connections to AI. A scaling law is how labs decide how big a model to train. Building one used to mean a dense grid, most of which is thrown away. Jha and colleagues (4 September) treat that collection as Bayesian optimization, keep the best-loss frontier, and fill the missing cells with the surrogate's guesses. They report fits close to the full grid at 10 to 100 times less compute.