

Dave's Daily

AI • Bayesian optimization US equities

tldr


1. OpenAI lost four safety researchers in a week. It fired three, and the man who wrote its safety reports quit, calling its culture “broken.”
2. Big Tech’s AI bond sales went from \$17 billion in 2024 to \$194 billion in the first half of 2026. The 10-year Treasury yield just hit a 24-year high, partly from that competition for capital.
3. AI hedge funds’ early edge faded to nothing after 2017, an NBER study finds. But they are far less correlated with each other than other funds are, which cuts against the herding fear.

AI

OpenAI’s safety bench empties. On Thursday OpenAI confirmed it had fired three safety and alignment researchers, Jasmine Wang, Tomek Korbak, and Mikita Balesni, for mishandling “sensitive company information.” The [Wall Street Journal](#) reports they allegedly shared it with an outside AI-safety group ([AFP](#)). On Saturday, David Robinson, who led the writing of OpenAI’s system cards, resigned in an Atlantic essay titled “I Quit OpenAI Because Its Culture Is Broken” ([Guardian](#)). His case: “iterative deployment” guarantees periodic failures, and their scale grows with capability. He wants aviation-style redundant safeguards and pressure from outside the company ([TechCrunch](#)). OpenAI says it pauses training or holds models back when needed.

Copying the reasoning. OpenAI says a coordinated campaign tried to extract its models’ hidden reasoning, peaking at 16,000 requests from over 4,000 users in two days in July. It ties a core cluster to people associated with Moonshot AI, maker of Kimi ([OpenAI](#); [CNBC](#)). No encryption was broken. Operators copied encrypted reasoning into a new chat and asked the model to transcribe it. OpenAI does not claim the material trained a Kimi model.

AI now borrows like a utility



People: a face that fools half of us. Tavus says 26 of 54 people who had a one-minute video call with its Griffin-Lite model believed their partner was human. Its previous system fooled 1 of 41 ([Tavus](#); [ThePrint](#)). The study is small, short, and run by the company. Tavus is holding the model back while it works on disclosure.

Souls and safety. Sam Altman said Saturday he is “very uncomfortable” with ascribing “religious force” to AI models, a jab at Anthropic co-founder Chris Olah’s talks with religious leaders about model consciousness ([Axios](#)). Hypothesis: the safety debate is splitting along two lines, catastrophic misuse and the moral status of models, and the labs no longer share a vocabulary for either.

Markets

AI is competing with Treasuries for money. The 10-year Treasury yield touched 5.34% on Thursday, its highest in 24 years. Reuters lists “increased competition for capital as companies issue debt to fund AI expansions” among the causes, alongside growth, energy prices, and last month’s Fed hike ([Reuters](#) [via](#) [Kitco](#)). J.P. Morgan finds tech credit spreads now trade wider than the investment-grade index, reversing the sector’s old premium. Oversubscription on the first surprise hyperscaler deal of the second half fell to 1.4 times, from 3.5–4.5 times earlier in the year (chart source). The mechanism is simple: every AI data center financed with bonds is a bond someone does not buy from the Treasury.

Does AI make quant funds herd? Not so far, says an NBER study of 7,896 US hedge funds from 2006 to 2024 ([NBER w35273](#); [NBER Digest](#)). AI-driven funds beat others by about 50 basis points a month before 2017, but the edge then became statistically indistinguishable from zero, even for early adopters. Pairwise return correlation among AI funds was 0.021, against 0.124 for non-AI funds. Only 89 funds qualified as AI-driven, so the sample is thin.

A language model that fakes order-book insight. Columbia’s Junxiao Chen and Paul Glasserman trained a language model on synthetic limit-order-book events ([arXiv:2608.23706](#)). It produced valid event sequences almost perfectly but never learned the state of the book, and it found predictability where none existed. For HFT desks testing generative order-flow models, valid output is not proof of a correct model.

Research paid for through execution. Citadel Securities took a minority stake in Wolfe Research on 2 October. Wolfe clients can now pay for its research through Citadel’s execution ([press release](#); [Investing.com](#)). A market maker now sells a bank’s research-plus-trading bundle without building a research desk.

Key numbers (this edition)

Item	Figure
10-year Treasury, 1 Oct	5.34%
Hyperscaler bonds, 2026 H1	\$194 billion
Extraction requests, 2 days	16,000
Griffin judged human	26 of 54
AI / non-AI fund correlation	0.021 / 0.124

Bayesian optimization

Data efficiency, and the cost of each decision. IQS-BO, posted 1 October by Luca Geminiani and Nadja Klein of KIT, replaces the acquisition function with a transformer ([arXiv:2610.01269](#)). Trained on synthetic functions, where the best candidate is always known, it learns the probability that each candidate is the maximizer. Picking the next query is then one forward pass. It matched or beat Gaussian-process BO on synthetic and real benchmarks while proposing queries about 11 times faster than a GP baseline, and 38 to 68 times faster for large batches.

Scaling. Donney Fan, Henry Moss, Jacob Gardner, Geoff Pleiss, and colleagues make latent-space BO fast enough for generate-and-screen pipelines ([arXiv:2609.19476](#)). A linear model on the sphere where high-dimensional latents concentrate gives nearly closed-form surrogates and acquisitions. It is over 100 times faster than baselines, in latent spaces up to 16,384 dimensions, with matching or better sample efficiency.

Connections to AI. Jeff Schneider’s group at Carnegie Mellon uses language models to feed BO what plain function values miss: training curves, expert notes, priors on where optima lie ([arXiv:2609.19437](#)). Their methods beat standard and LLM-based optimizers on hyperparameter benchmarks and a nuclear-fusion control task. Put together, these papers show the field moving the expensive judgment from the evaluation loop into a pre-trained model.