

Dave's Daily

AI • Bayesian optimization • US equities

tldr

1. Anthropic will send opt-in maintainers bug reports that no person has read. Of 97 high and critical findings its testers did review, 85 cleared the human-disclosure bar, 11 were real duplicates, and one was invalid.
2. OpenAI told investors September annualized revenue was about \$50 billion. A widely cited \$68 billion had included partners' gross revenue. Several chip and cloud shares fell the same day.
3. A new optimizer searches a map of events. With the mesh fixed, the expensive refit does not grow as events arrive. On a one-dimensional test, simple regret reached zero by iteration 15.

AI

Reports before a person reads them.

On Thursday Anthropic opened [OSS Scanner](#). Projects that opt in get periodic scans from its strongest models, including Claude Mythos, at no charge. The write-up is the model's: a reproducer, an explanation, and a patch when one can be written. No person triages it first. Maintainers enroll with a pull request. Anthropic says it will favor projects with a critical effect on infrastructure, case by case, on a standard it compares to Google's OSS-Fuzz.


The queue is the constraint. The lab says six months of scanning produced more than 29,000 candidate vulnerabilities, of which people reviewed about 6,000. It has already forwarded nearly 5,000 unreviewed reports to maintainers who asked for the pile. Human-checked disclosures continue for projects that cannot take the raw stream.

The chart is the audit Anthropic published. Testers who already review its coordinated disclosures looked at 97 critical and high findings from an early scanner, across 48 projects. Eighty-five met that bar. Eleven were real but duplicated a known issue or another hit in the scan. One was invalid. Anthropic says maintainers have seldom rejected a high or critical finding since, though some call the severity high or say the model missed the threat model.

wolfSSL's Todd Ouska, quoted there, said 74 reports left two invalid and five that became CVEs. That is one project.

Anthropic also says models on the CyberGym benchmark found under 20% of the bugs at the start of last year and over 85% this year. Hypothesis: an opt-in maintainer is trading reading time, including false alarms, for a head start on attackers who can now draft an exploit quickly.

A person checked 97 high and critical findings



Gold is the count. The pale track is all 97. The axis starts at zero. These are high and critical findings from an early scanner, reviewed by the penetration testers who also check Anthropic's human disclosures. Source: [Anthropic](#), 8 October 2026.

A coworker with an inbox.

Thursday morning Google Cloud showed a Gemini agent at Gemini at Work, in private preview, with no general-release date ([Google](#); [VentureBeat](#)). A spokesperson told VentureBeat it can work for hours or days, keep one identity across devices, call Claude as well as Gemini, and will cost nothing extra where Gemini Enterprise is already paid for. In the coworker setup Google describes an agent with its own email, calendar, Drive, and directory listing. Google has not said whether every agent gets that account, whether a license is extra, or whether a prompt sent to Claude leaves Google's boundary. It says actions are logged and an agent sees only files shared with it. It has not published a test of how often a multi-day job finishes correctly.

Hypothesis: a mailbox that colleagues treat as a person is a new place for a bad send, even when sharing rules are ordinary.

The firing, in their words.

Jasmine Wang, Tomek Korbak, and Mikita Balesni, dismissed by OpenAI last week, published an open letter Thursday denying they had mishandled information outside their jobs ([TechCrunch](#)). Wang wrote that the cited email was recruiting access IT had not removed, opened by mistake and reported within minutes. A memo shared with TechCrunch says they were not fired for raising safety concerns. A spokesperson cited a "pattern of misconduct" and did not name the rule. The letter and the company disagree. What is public is that disagreement.

Markets

Whose dollars were in the \$68 billion.

OpenAI told investors that annualized revenue at the end of September was about \$50 billion, CNBC confirmed Thursday, under the roughly \$68 billion figure reported late last month ([CNBC](#); the Financial Times had the \$50 billion first). A person familiar with the presentation said the higher figure included partners' gross revenue, so investors could line OpenAI up against Anthropic. The same person said the third quarter's total run rate grew 77% and the enterprise run rate 107%. CNBC put Thursday's moves next to that clarification: Nvidia about 3% lower, Oracle nearly 6%, CoreWeave nearly 8%. The closes are not the story. A private company changed what the number includes, and firms that sell chips, cloud, and data centers moved with it. Hypothesis: this trade is priced off a slide, so the definition of revenue is part of the risk.

State the intent, then accept the order.

Fortune reported Thursday that Sequoia led a \$30 million seed round for Catalyst, whose agents turn a sentence into a strategy and then an order ([Fortune](#)).

Other investors include Jump Trading, Peak XV, Lux Capital, AntiFund, Coinbase, and Premji Invest. Cofounder Dylan Iskandar said a person still has to accept, and that the aim is for the person to state an intent and leave the rest to the agents.

Catalyst says a pilot with active traders did hundreds of millions of dollars of volume in a couple of weeks. That figure is the company's. The product was still waitlisted Thursday. This is not a high-frequency desk. It is a retail path from a sentence to a ticket, with the accept button as the control and a proprietary trading firm on the cap table. Hypothesis: the accept step becomes optional once the screen is built so that refusal is the slow choice.

Bayesian optimization

Scaling: fit the field on a fixed mesh.

Standard Bayesian optimization gets a number from each trial and fits a Gaussian process whose dominant cost grows with the cube of the trial count. A log-Gaussian Cox process is the analogous model when the data are events in space or time, not a score. Kaichuang Yang, Harvard Rue, and Jakob Zeitler (KAUST and Oxford) optimize that intensity ([arXiv:2610.12213](#), 8 October). They call it ISBO. A Matérn field lives on a mesh of triangles, so the latent weights are a sparse Gaussian Markov random field. Modern integrated nested Laplace approximation keeps that latent vector on the mesh. Their theorem: if the mesh and the model stay fixed, the sparse factorization does not grow as more events appear. Building the likelihood still can. On Porto taxi traces the refit stays nearly flat out to 55,000 events, while two reduced-rank approximations slow down as the count rises. On 342 firearm reports in Washington, D.C., May–June 2022, one fit took 2.27 seconds, against 35.74 seconds for an earlier kernel method.

The mechanism.

A query opens a region and returns the events inside it. The acquisition is an upper confidence bound on the log-intensity, with an exploration bonus that shrinks, plus a mask so a seen region is not opened again. Penalized-complexity priors keep the surface simple while few events are in hand. The paper links an anonymous code repository. The durable pointer is the arXiv record.

Data efficiency.

On a known one-dimensional intensity, ten seeds, mean simple regret hit zero by iterations 14–15. A fixed exploration bonus was slower and more variable. Dropping the penalized-complexity prior left a small regret and a visible band. The earlier kernel baseline stayed higher, and on the same temporal task its simple regret leveled off while ISBO's fell toward zero. There is no classical regret proof. The claim is empirical: find the busy region with few revealed windows.

Connections to AI.

The introduction names financial events, disease cases, and neural spikes as this kind of pattern. The speed tests are crime counts and taxi traces, not an order book. Hypothesis: what would carry into trading is the flat refit on a fixed mesh, when the observation is a cloud of events rather than one score. Yesterday's paper spent idle time on a better average of a score. This one changes the observation.

Updated Friday 9 October 2026, before the US open. The same text ships as a phone page and a letter page. OSS Scanner, the Gemini agent, the open letter, the revenue account, and the Catalyst round are from Thursday 8 October. The optimization paper is from 8 October. Market notes describe an accounting basis and a product, not a forecast. Created with [malvin](#).