Dave's Daily

tldr

1. Anthropic told investors on 14 October that third-quarter revenue was $16.8 billion, from $11.6 billion in the second quarter and $4.7 billion in the first, and that the quarter was profitable on an adjusted basis. **2.** Bloomberg is opening ASKB, its Terminal chat, to every subscriber and letting them add their own documents. More than a third of subscribers already use it. **3.** A Munich group searches for low tail risk when the Gaussian-process kernel is unknown. On a 26-setting SAT solver, the other methods they tried finished above a CVaR regret of 50.

AI

On 14 October in San Francisco, Anthropic told prospective investors that third-quarter revenue rose 45 percent, to $16.8 billion, from $11.6 billion in the second quarter and $4.7 billion in the first. Two people who attended told Reuters. The Financial Times reported the third-quarter figure first. Anthropic also told the group the quarter was profitable on an adjusted basis. Anthropic did not answer Reuters's request for comment.

The chart is those three quarters on one scale. The rise from the first quarter to the second is $6.9 billion. The rise from the second to the third is $5.2 billion. The second quarter was about 2.5 times the first. The third is 1.45 times the second.

The meeting was a preview of a listing. Reuters reports that Anthropic is targeting a valuation above $2 trillion, and that the debut is likely to wait until after the November midterm elections. Last year revenue rose twelvefold to nearly $4.6 billion, while the operating loss was more than $8 billion, excluding writedowns mostly tied to earlier fundraising. The company has committed at least $518 billion over a decade for cloud services, computing, and infrastructure. A prospectus Reuters saw last month warned that advanced AI could pose catastrophic or existential risks to humanity. At the Wednesday meeting Amodei talked about safety. In August, Reuters reported that an annualized run rate had topped $65 billion by the end of July. A finished quarter of $16.8 billion, repeated four times, is $67.2 billion. A run rate and a finished quarter are different measures, and here they sit close together.

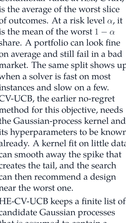

Anthropic quarterly revenue in 2026, in billions of dollars. Two people who attended the 14 October investor day gave Reuters these figures. The third quarter is 45 percent above the second. The axis starts at zero and ends at 20. Source: Reuters, 14 October 2026.

Markets

On 14 October, the Wall Street Journal reported that Bloomberg is making ASKB available to every Terminal subscriber and letting them place their own content next to the financial information already in the workflow. Bloomberg began a limited rollout early this year. The Journal says more than one-third of Terminal subscribers now use ASKB. The interface sits on a system of models and agents.

A subscriber can ask in ordinary language and attach the firm's own documents, so an answer can draw on private material and on Bloomberg's data, news, and analytics together. The Journal describes a research desk putting its files into the same conversation as the market record, on the terminal US equity desks already pay for.

Bayesian optimization

arXiv:2610.18306, posted 13 October, is *Bayesian Optimization of CVaR with Unknown Hyperparameters*, by Valentin Margraf, Jonas Hanselle, Viktor Bengs, Serafina Lebedeva, Eyke Hüllermeier, and Marcel Wever, at LMU Munich, the Munich Center for Machine Learning, IDFKI, RPTU Kaiserslautern-Landau, Leibniz University Hannover, and the L3S Research Center.

Conditional value at risk, CVaR, is the average of the worst slice of outcomes. At a level α, it is the mean of the worst $1-\alpha$ share. A portfolio can look fine on average and still fail in a bad market. The same split shows up when a solver is fast on most instances and slow on a few. CV-UCB, the earlier no-regret method for this objective, needs the Gaussian-process kernel and its hyperparameters to be known already. A kernel fit on little data can smooth away the spike that creates the tail, and the search can then recommend a design near the worst one.

HE-CV-UCB keeps a finite list of candidate Gaussian processes that is assumed to contain a correct one. Each round it picks the design and the candidate with the most optimistic lower bound on CVaR, and it scores that design with the most pessimistic candidate still in the list. It then measures how the environment where those two disagree the most. A candidate whose residuals exceed its own confidence bound is removed. With high probability the correct candidate stays. Cumulative regret matches CV-UCB's oracle rate up to a factor equal to the number of candidates. The returned design, among those tried, is the one with the smallest worst-case posterior-mean CVaR across whoever remains. The search is optimistic. The recommendation is conservative.

On a toy function drawn from a known process, the pool held that process and two oversmoothed ones. HE-CV-UCB dropped the oversmoothed models after a median of 4 and 13 iterations and kept true one on every seed. The regret of its recommended design matched an oracle that was handed the true kernel. On a spike function, methods that pick or average models by likelihood plateau near a CVaR regret of 1. HE-CV-UCB reaches near zero, because the candidates already agree on the bulk, so the next query lands on the spike.

The larger tests transfer one Gaussian process from each related source problem. With a robot pushing toward a shifted goal, HE-CV-UCB matches the best baseline, CV-TS. With a heat-releasing reactor scaled up from smaller vessels, models fit on the small reactors underestimate the failure tail, and HE-CV-UCB finishes with the lowest CVaR regret.

Configuring the Spear SAT solver is the wide search: 26 settings, 604 software-verification instances, runtimes censored at 300 seconds, and eight source clusters with one held out. The other methods finish with CVaR regret above 50. HE-CV-UCB is lower. The guarantee needs a correct model somewhere in the pool. On the transfer tasks the authors cannot promise that, and they say so. The code path is BoTorch.

Updated 15 October 2026. Phone and letter match. Sources: Reuters, 14 October; Wall Street Journal, 14 October; arXiv, 13 October. The revenue figures are what two attendees told Reuters. Created with malvin.