Dave's Daily

tldr

1. On 18 October Anthropic said it will commit 5 percent of its shares to a new public-benefit fund, subject to stockholder approval. The board may add 2 percentage points for every 50 percent rise in company value above a baseline Anthropic has not published.
2. On Saturday, with U.S. stock markets closed, about $51.6 million of leveraged chip-stock contracts were forced-closed in ten minutes. Orders on 18 mance at 21:21:59 UTC reached Hyperliquid less than a second later.
3. A 16 October paper keeps a molecule generator frozen and reweights its samples with expected improvement. After 300 scored molecules, the flow variant's mean on JNK3 was 0.778, against 0.412 for the next method.

AI

On 18 October Anthropic said it is committing shares to the Anthropic Public Benefit Fund. Subject to stockholder approval, 5 percent of the shares will go now. The board may add another 2 percent for every 50 percent increase in company value above a set baseline. Anthropic said it will publish those valuation milestones in the coming months. The fund is a wholly owned LLC. Anthropic says that form lets it move between gifts to nonprofits and direct distributions to particular groups. None of the equity the LLC holds will be sold by lobbying or spending, in the United States or abroad.

The example Anthropic gives is the labor market. It says AI may create large enough gains that it may also unsettle many workers. The company has already committed $200 million to studying and addressing that issue, and it says the fund can take up larger responses. The co-founders plan to start donating to donate 80 percent of their own equity, and about half of employees take part in a match of pledged equity. With Sunday's commitment, Anthropic puts the public-good share at roughly 20 percent, plus any further points if the milestones are met. The BBC reported the announcement the same day.

Share of Anthropic the fund may hold. Longer is a larger share. The area starts at zero and ends at 16. The baseline bar is the 5 percent that needs shareholder approval. Each later bar adds 2 points, and only if the board chooses, for another 50 percent rise above the last one the company has not published. Source: Anthropic, 18 October 2026.

The same day, Microsoft left Nscale's planned datacentre in Loughton, Essex. The Guardian reports that ministers had called the site the largest sovereign AI datacentre in the UK, with an opening planned for next year. UK Power Networks has told Nscale it cannot supply enough power until the site's grid is mid-2030s. Nscale said it is looking at power generated on the site, and that Microsoft's contracts with other Microsoft contracts in place.

Markets

On Saturday the shares were not trading, and the contracts that track them were. Reporting Malwa, writing in CoinDesk on 18 October, reports that about $51.6 million of perpetual futures tied to chip stocks and semiconductor funds were force-closed between 21:22 and 21:31:59 UTC. Less than a second later, orders arrived on Hyperliquid, where Trade XYZ lists these contracts. About ten additional venues had against six chip-linked contracts within 1.8 seconds. In that first minute Binance Futures had about $21.9 million of SanDisk trading, against $3.1 million on Trade XYZ.

These contracts follow a share and keep trading through the weekend. With the exchange closed, there is no stock print to follow. CoinDesk reports that Trade XYZ's rules let trades on that platform set its reference prices until cooled-down rules re-start at the U.S. open. The ten minutes show how forced selling moves between venues while the shares sit still.

Bayesian optimization

arXiv:2610.22763, submitted 16 October, is *Acquisition Function-Guided Generation for Bayesian Optimization over Discrete Spaces*, by Tommaso Ancilli and Thomas Wedenig (TU Graz), Thierry Moss (Lorraine), and Robert Peharz (TU Graz). The code release is anonymous.anonymous.science/r/afgg.

The design a seed molecule as a molecule written as a SELFIES string of fixed length. Valid strings are ignored, and, and the search has no gradient. A pretrained generator already concentrates on strings that look like molecules, and the Bayesian step it frozen. The score model is a Gaussian process on the molecule, with a Tanimoto kernel on Morgan fingerprints, usually 2,048 bits. Expected improvement ranks a molecule by how far it might beat the best score so far. The next draw comes from the generator, reweighted by that rank. A difference turns the reweighted target into a smooth average, which gives the sampler a gradient. After each new score, the Gaussian process and the acquisition update. The generator stays as it was.

Two generators are tested. CDCD is a categorical diffusion model the authors trained. SeqFlow is a pretrained autoregressive flow from Lee and coauthors (2025), used unchanged. AFGG needs a generator that can be guided and an acquisition that is non-negative and can be scored one molecule at a time. Expected improvement qualifies. Each draw starts near the incumbent, because the authors report that sample quality falls when a high-dimensional search start sits off the thin shell where that Gaussian concentrates.

The sequential benchmark is the Practical Molecular Optimization suite: 25 objectives, 10 starting scores, then 300 more, one per round, as a mean and a sample standard deviation over five seeds. A count of Table 6 shows the SeqFlow mean is highest on one of the eleven baselines on 18 of the 25 tasks. On 11 of the 25, the lead also exceeds SeqFlow's own sample standard deviation, the comparison the paper uses when it shades a cell.

On JNK3, SeqFlow finishes at 0.778 (standard deviation 0.038) and NFBO at 0.412 (0.095). Celecoxib rediscovery is 0.836 (0.125) against 0.551 (0.085). Albuterol similarity is 0.975 (0.038) against InvBO at 0.629 (0.079). The diffusion variant, CDCD, is often level with the stronger baselines, and on fexofenadine it leads SeqFlow, 0.763 (0.017) against 0.741 (0.016).

Valsartan SMARTS ends at 0.000 for SeqFlow, CDCD, and the baselines; NFBO's mean is 0.015 (0.034). On logP, CoBO is 65.866 (10.720) and SeqFlow is 35.734 (10.360). The paper says SeqFlow and synthetic accessibility reward large or hard-to-make molecules, a departure from the training distribution, and the latent-space methods lead there. A per-position inner sample took 0.36 hours per seed against 8.77 hours for a 32-step rollout. On six objectives and three seeds, the paper repeats the protocol and the seed spread at the slower setting. In rounds of five, separate tits and one-standard-deviation bars.

Updated 19 October 2026. Phone and letter match. Sources: Anthropic and the BBC, 18 October; the Guardian, 18 October; CoinDesk, 18 October, on a Saturday session; arXiv, 16 October 2026. Code: anonymous.anonymous.science/r/afgg, as released. The liquidation total is the sum of forced contracts, not a forecast. Created with malvin.